

Grace Chang

(224) 201-8256 · chag60460@gmail.com · chag60460.github.io/grace-portfolio

WORK EXPERIENCES

AI Software Engineer · *Microsoft (Experiences & Devices)*

Nov 2024 – Present

- Led end-to-end redesign and development of Windows Privacy Settings page that serves 4M+ monthly active devices, improving navigation efficiency and feature discoverability by 10% through gathering telemetry data, designing architectural solutions, leading specification evaluations, and implementing new interface that resulted in a 1.57 second boost in render time.
- Prototyped and integrated agentic frameworks to capture, triage, root-cause, and mitigate crashing bugs with 50k daily hits from retail users, reducing manual synchronization across 100+ Settings page entries and 7 partner services while automating guardrails to enable early detection of any regression.
- Pioneered and built LLM-based orchestrator adopted across the entire engineering organization to standardize agentic workflows and turn tribal knowledge into shared AI capability backed modules, reducing new hire onboarding friction by 87.5% through chaining MCP tool calls to surface product overviews, provision permissions and repositories, and automate all processes across the software development lifecycle.

Distributed Software Engineer · *Microsoft (Azure Core)*

Aug 2023 – Nov 2024

- Optimized distributed storage platform serving 40B+ daily requests, reducing operational latency up to 250% at P99.9 via refining software infrastructure, validating storage synchronization, and managing traffic load.
- Integrated Azure Throttling Service into Azure Resource Manager to help the throttling service attain competitive market rate of 700k requests per hour, through enabling quick configuration of size limits and implementing emergency configuration service that support dynamic adjustment based on real-time system load across 20,000 VM admissions per minute with 0% READ throttling or WRITE throttling.
- Isolated Cosmos database containers optimized for queries, helping reduce CPU usage from an average of 70% to 3% and decrease time duration of resources queries from 330 milliseconds to 5 milliseconds by implementing a flexible query mode based on resource availability and container maximization.

PROJECTS

Slack Agent — *LLM-Powered Slack Agent*

- Implemented an LLM-powered Slack agent that automates a volunteer team's recurring meeting ops by parsing a natural-language message with Llama 3.1 8B (Groq) into a strict JSON schema and dispatching Google Calendar, Forms, and Gmail tool calls behind human-in-the-loop confirmation — replacing manual coordination across four tools.

AI Temperature Monitor — *Microsoft ImPaCT, Experiences & Devices*

- Built an end-to-end ML system pairing Arduino IoT sensor streams with a TensorFlow forecasting model for real-time temperature prediction and automated control.
- Designed the full data-ingestion, model-training, and inference pipeline with a feedback loop between physical actuation and model predictions.

Magdalene X² — *Multi-Agent Health-Intelligence Product*

- Architected an autonomous multi-agent health companion that unifies women's labs, cycle, nutrition, fitness, and sleep — through user research, synthetic reviews, PRD, iterative prototypes, and code development — to surface a cross-domain insight and make it actionable via a safe closed-loop N-of-1 experiment, a gap none of 30+ analyzed products fill.

SKILLS

AI & Systems	Prompt Engineering, Tool Calling & Function Dispatch, LLMs, Agentic System Design, MCP / Tool Orchestration, Model Selection & Benchmarking, TensorFlow
Software Development	Python, C#, C++, Java, Flask, MongoDB, MySQL, Git, TypeScript, JavaScript, HTML, CSS, React, Vue, Node.js
Prototyping & Design	Figma, Unity, TouchDesigner, Arduino

EDUCATION

Massachusetts Institute of Technology · Cambridge, MA

Artificial Intelligence · Tangible User Interface · Algorithms · Software Construction · Software Studio

Wellesley College · Wellesley, MA

Databases with Web Interface · Human Computer Interaction